

# COL331 Operating Systems Assignment 3

Abhinav Rajesh Shripad      Entry Number: 2022CS11596

Jahnabi Roy      Entry Number: 2022CS11094

April 30, 2025

## Memory Printer in xv6

### 0.1 Ctrl + I Implementation

- For this implementation, like in the last assignment changes were made to the `console.c` file for recognising the keyboard interrupt.
- Changes were made to the `proc.c`. The `ctrl_memoryprint` function prints the detected keys and prints PID and the `NUM_PAGES`.

The main relevant code sections are as follows:

In `console.c`:

```
1 void consoleintr(int (*getc)(void))
2 {
3     int to_memoryprint = 0;
4     (...)
5     case C('I'):
6         to_memoryprint = 1;
7         break;
8     (...)
9     if (to_memoryprint) {
10         ctrl_memoryprint();
11     }
12     (...)
13 }
```

In `proc.c`:

```
1 void ctrl_memoryprint(void)
2 {
3     struct proc *p;
4     int num_pages;
5     acquire(&ptable.lock);
6     cprintf("\nCtrl+I is detected by xv6\n");
7     cprintf("PID NUM_PAGES\n");
8     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
9     {
10         if (p->pid > 0)
11         {
12             num_pages = PGROUNDUP(p->sz) / PGSIZE;
13             cprintf("%d %d\n", p->pid, num_pages);
14         }
15     }
16     release(&ptable.lock);
17 }
```

## Page Swapping in xv6

### 0.2 Implementation Details

- System calls, `getrss` and `getNumFreePages` were implemented as according to implementation done for assignment 1.

- `find_victim_process` defined in `proc.c`, selects a victim process — i.e., one from which a page can be evicted — based on highest resident set size (RSS). `find_victim_page` defined in `proc.c` implements a Second-Chance (Clock)-style page replacement policy: it tries to find a page that hasn't been accessed recently. If all pages are accessed, give some of them a "second chance" by clearing `PTE_A`, then try again.
- Changes made in `kalloc.c`, `exec.c`, `mkfs.c`, `fs.c` and `trap.c` to implement the copy-on-write mechanism.
- This implementation adds dynamic swapping to xv6 by monitoring free physical memory and triggering page swaps when it falls below a threshold. The `kalloc()` function checks the number of free pages and, if low, calls `swapOut()` to free up space by evicting a page. `swapOut()` selects a victim process with the highest memory usage ('rss') using `find_victim_process()`, and within that process, it finds a suitable page to swap out using `find_victim_page()`, preferring cold (unaccessed) pages. The selected page is written to disk, its metadata is updated, and its physical memory is freed. This implements an adaptive, LRU-like swapping mechanism to maintain system responsiveness under memory pressure.
- `pageswap.c` was created and the code for the same is shown below.

The main relevant code is as follows:

In `sysproc.c`:

```

1  (...)
2  int sys_getNumFreePages(void)
3  {
4      return num_of_FreePages();
5  }
6
7  int sys_gettrss()
8  {
9      print_rss();
10     return 0;
11 }
12 (...)

```

In `proc.c`:

```

1  struct proc *find_victim_process(void)
2  {
3      struct proc *p;
4      struct proc *victim_p = NULL;
5      int pid = 100000;
6      uint highest_rss = 0;
7
8      acquire(&ptable.lock);
9      for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
10     {
11         if (p->pid < 2)
12             continue;
13
14         if (p->rss >= highest_rss)
15         {
16             if (p->pid < pid)
17             {
18                 pid = p->pid;
19             }
20
21             victim_p = p;
22             highest_rss = p->rss;
23         }
24     }
25     release(&ptable.lock);
26
27     return victim_p;
28 }
29
30 uint find_victim_page(struct proc *p)
31 {
32     uint i;

```

```

33     uint sz = p->sz;
34     pte_t *pte;
35     uint count = 0;
36
37     for (i = 0; i < sz; i += PGSIZE)
38     {
39         pte = walkpgdir(p->pgdir, (void *)i, 0);
40         if ((*pte & PTE_P) && !(*pte & PTE_A) && (*pte & PTE_U))
41         {
42             p->rss -= PGSIZE;
43             P2V(PTE_ADDR(*pte));
44             return i;
45         }
46         if (*pte & PTE_P)
47             count++;
48     }
49
50     count = (count / 10) + 1;
51
52     for (i = 0; i < sz; i += PGSIZE)
53     {
54         pte = walkpgdir(p->pgdir, (void *)i, 0);
55         if ((*pte & PTE_P) && (*pte & PTE_A) && (*pte & PTE_U))
56         {
57             *pte &= ~PTE_A;
58             count--;
59         }
60         if (!count)
61             break;
62     }
63     return find_victim_page(p);
64 }

```

In exec.c:

```

1     (...)
2     curproc->rss = curproc->sz;
3     (...)

```

In kalloc.c

```

1     (...)
2
3     struct
4     {
5         struct spinlock lock;
6         int use_lock;
7         uint num_free_pages;
8         struct run *freelist;
9     } kmem;
10
11     struct
12     {
13         int use_lock;
14         struct spinlock lock;
15         int ref[PHYSTOP / PGSIZE];
16     } rmap;
17
18     int getRmapRef(uint pa)
19     {
20         if (rmap.use_lock)
21             acquire(&rmap.lock);
22         int num = rmap.ref[pa / PGSIZE];
23         if (rmap.use_lock)
24             release(&rmap.lock);
25         return num;
26     }
27
28     void setRmapRef(uint pa, int val)
29     {
30         if (rmap.use_lock)
31             acquire(&rmap.lock);
32         rmap.ref[pa / PGSIZE] = val;
33         if (rmap.use_lock)

```

```

34         release(&rmap.lock);
35     }
36
37     void incRmapRef(uint pa)
38     {
39         if (rmap.use_lock)
40             acquire(&rmap.lock);
41         ++rmap.ref[pa / PGSIZE];
42         if (rmap.use_lock)
43             release(&rmap.lock);
44     }
45
46     void decRmapRef(uint pa)
47     {
48         if (rmap.use_lock)
49             acquire(&rmap.lock);
50         --rmap.ref[pa / PGSIZE];
51         if (rmap.use_lock)
52             release(&rmap.lock);
53     }
54
55     void kinit1(void *vstart, void *vend)
56     {
57         (...)
58         initlock(&rmap.lock, "rmap");
59         (...)
60         rmap.use_lock = 0;
61         (...)
62     }
63
64     void kinit2(void *vstart, void *vend)
65     {
66         (...)
67         rmap.use_lock = 1;
68         (...)
69     }
70
71     (...)
72
73     void kfree(char *v)
74     {
75         (...)
76         decRmapRef(V2P(v));
77         if (getRmapRef(V2P(v)) > 0)
78         {
79             return;
80         }
81
82         (...)
83         kmem.num_free_pages += 1;
84         (...)
85     }
86
87     int Th = THRESHOLD;
88     int Npg = NPG;
89
90     char *
91     kalloc(void)
92     {
93         struct run *r;
94         int free = kmem.num_free_pages;
95
96         if (free <= Th)
97         {
98             cprintf("Current Threshold = %d, Swapping %d pages\n", Th, Npg);
99             for (int i = 0; i < Npg; i++)
100             {
101                 swapOut();
102             }
103             // Th = Th * (100 - BETA) / 100;
104             // Npg = Npg * (100 + ALPHA) / 100;
105             Th -= (Th * BETA) / 100;
106             Npg += (Npg * ALPHA) / 100;

```

```

107         if (Npg > THRESHOLD)
108             Npg = THRESHOLD;
109     }
110
111     if (kmem.use_lock)
112         acquire(&kmem.lock);
113     r = kmem.freelist;
114     if (r)
115     {
116         kmem.freelist = r->next;
117         kmem.num_free_pages -= 1;
118         setRmapRef(V2P(r), 1);
119     }
120
121     if (kmem.use_lock && 1)
122         release(&kmem.lock);
123
124     if (SWAPON)
125     {
126         if (!r)
127         {
128             swapOut();
129             return kalloc();
130         }
131     }
132
133     return (char *)r;
134 }
135
136 uint num_of_FreePages(void)
137 {
138     acquire(&kmem.lock);
139     uint num_free_pages = kmem.num_free_pages;
140     release(&kmem.lock);
141     return num_free_pages;
142 }

```

In mkfs.c

```

1  (...)
2  nmeta = 2 + SWAPBLOCKS + nlog + ninodeblocks + nbitmap;
3  nblocks = FSSIZE - nmeta;
4  sb.size = xint(FSSIZE);
5  sb.nblocks = xint(nblocks);
6  sb.ninodes = xint(NINODES);
7  sb.nlog = xint(nlog);
8  // sb.logstart = xint(2);
9  sb.logstart = xint(SWAPBLOCKS + 2);
10 // sb.inodestart = xint(2+nlog);
11 sb.inodestart = xint(2 + SWAPBLOCKS + nlog);
12 // sb.bmapstart = xint(2+nlog+ninodeblocks);
13 sb.bmapstart = xint(2 + SWAPBLOCKS + nlog + ninodeblocks);
14 sb.swapblocks = xint(SWAPBLOCKS);
15 sb.swapstart = xint(2);
16
17 printf("nmeta %d (boot, super, swap blocks %u log blocks %u inode blocks %u, bitmap
18       blocks %u) blocks %d total %d\n",
19       nmeta, SWAPBLOCKS, nlog, ninodeblocks, nbitmap, nblocks, FSSIZE);
20
21 (...)
22 void balloc(int used)
23 {
24     uchar buf[BSIZE];
25     int i, j;
26
27     printf("balloc: first %d blocks have been allocated\n", used);
28
29     for (j = 0;; j++)
30     {
31         int start = j * BSIZE * 8;
32         if (start >= used)
33             break;
34         int end = start + BSIZE * 8;

```

```

35     int chunk = end - start;
36     if (chunk > used - start)
37         chunk = used - start;
38
39     // Read the j-th bitmap block (sector 2 + j)
40     rsect(2 + j, buf);
41     bzero(buf, BSIZE); // Initialize all bits to 0
42
43     for (i = 0; i < chunk; i++)
44     {
45         buf[i / 8] |= 0x1 << (i % 8);
46     }
47
48     wsect(2 + j, buf);
49 }
50 }

```

In fs.c

```

1     (...)
2     cprintf("sb: size %d swapblocks %d nblocks %d ninodes %d nlog %d logstart %d\
3     inodestart %d bmap start %d\n", sb.size, SWAPBLOCKS, sb.nblocks,
4     sb.ninodes, sb.nlog, sb.logstart, sb.inodestart,
5     sb.bmapstart);
6     (...)

```

In trap.c

```

1     (...)
2     case T_PGFLT:
3         page_fault_handler();
4         break;
5     (...)

```

In pageswap.c:

```

1     (...)
2     #define SWAPPAGES (SWAPBLOCKS / 8)
3
4     struct swapinfo
5     {
6         int pid;
7         int page_perm;
8         int is_free;
9     };
10
11     struct swapinfo swp[SWAPPAGES];
12
13     // buffer
14
15     void swapInit(void)
16     {
17         for (int i = 0; i < SWAPPAGES; i++)
18         {
19             swp[i].page_perm = 0;
20             swp[i].is_free = 1;
21         }
22     }
23
24     void swapOut()
25     {
26         struct superblock sb;
27         readsb(ROOTDEV, &sb);
28         // find victim page
29         struct proc *p = find_victim_process();
30         uint victim_page_VA = find_victim_page(p);
31
32         pte_t *victim_pte = walkpgdir(p->pgdir, (void *)victim_page_VA, 0);
33
34         int i;
35         // find free swap slot
36         for (i = 0; i < SWAPPAGES; ++i)
37         {
38             if (swp[i].is_free)

```

```

39     {
40         break;
41     }
42 }
43
44 swp[i].page_perm = 0;
45 swp[i].is_free = 0;
46 swp[i].pid = p->pid;
47
48 // update swap slot permissions to match victim page permissions
49 swp[i].page_perm = PTE_FLAGS(*victim_pte);
50 char *pa = (char *)P2V(PTE_ADDR(*victim_pte));
51 uint addressOffset;
52 for (int j = 0; j < 8; ++j)
53 {
54     addressOffset = PTE_ADDR(*victim_pte) + (j * BSIZE);
55     struct buf *b = bread(ROOTDEV, sb.swapstart + (8 * i) + j);
56     memmove(b->data, (void *)P2V(addressOffset), BSIZE);
57     b->blockno = (sb.swapstart + (8 * i) + j);
58     b->flags |= B_DIRTY;
59     b->dev = ROOTDEV;
60     bwrite(b);
61     brelse(b);
62 }
63 (*victim_pte) = ((sb.swapstart + (8 * i)) << 12) & (~0xFFF);
64 *victim_pte &= ~PTE_P;
65 lcr3(V2P(p->pgdir));
66 kfree(pa);
67 }
68
69 void swapIn(char *memory)
70 {
71     struct proc *p = myproc();
72     uint addr = rcr2();
73     pde_t *pd = p->pgdir;
74     pte_t *pg = walkpgdir(pd, (void *) (addr), 0);
75
76     uint swapSlot = (*pg >> 12); // physical address of swap block
77     int swapIdx = (swapSlot - 2) / 8;
78     *pg = (V2P((uint)memory) & (~0xFFF)) | swp[swapIdx].page_perm;
79     // int swapSlot = swap; // swap block number (convert it to integer)
80     for (int i = 0; i < 8; i++) // writes page into physical memory
81     {
82         struct buf *b = bread(ROOTDEV, swapSlot + i);
83         cprintf("\n");
84         memmove(memory, b->data, BSIZE);
85         brelse(b);
86         memory += BSIZE;
87     }
88     swp[swapIdx].is_free = 1;
89     swp[swapIdx].pid = 0;
90 }
91
92 void freeSwapSlot(int pid)
93 {
94     struct superblock sb;
95     readsb(ROOTDEV, &sb);
96     int i;
97     for (i = 0; i < SWAPPAGES; i++)
98     {
99         if (swp[i].pid == pid)
100             break;
101     }
102     for (int j = 0; j < 8; j++)
103     {
104         struct buf *b = bread(ROOTDEV, sb.swapstart + (8 * i) + j);
105         memset(b->data, 0, BSIZE);
106         brelse(b);
107     }
108     swp[i].is_free = 0;
109     swp[i].pid = 0;
110 }

```

### 0.3 Role of $\alpha$ and $\beta$ in efficiency

This analysis examines how the parameters  $\alpha$  (swap-out growth rate) and  $\beta$  (threshold reduction rate) impact the efficiency of the adaptive page replacement strategy. Using experimental data (as available in `swap_efficiency.csv`), we evaluate three key metrics:

- Swap Operations: Number of swap-out cycles
- Total Pages Swapped: Cumulative pages moved to disk
- Threshold Adaptation: Pattern of threshold ( $Th$ ) reduction

#### 0.3.1 Role of $\beta$ (Threshold Reduction Rate)

$$Th_{new} = \left\lfloor Th_{current} \cdot \left(1 - \frac{\beta}{100}\right) \right\rfloor \quad (1)$$

Table 1: Impact of  $\beta$  Values ( $\alpha = 10$ )

$\beta$	Swap Ops	Total Pages	Final $Th$	Behavior
10	31	124	9	Gradual decay, frequent swaps
50	17	68	1	Rapid decay, fewer swaps
100	16	64	0	Threshold collapses immediately

Key observations:

- High  $\beta$  (50-100): Reduces swap frequency but risks premature threshold collapse
- Low  $\beta$  (10-30): Maintains memory pressure awareness at cost of higher I/O

#### 0.3.2 Role of $\alpha$ (Swap-Out Growth Rate)

$$Npg_{new} = \min \left( LIMIT, \left\lfloor Npg_{current} \cdot \left(1 + \frac{\alpha}{100}\right) \right\rfloor \right) \quad (2)$$

Table 2: Impact of  $\alpha$  Values ( $\beta = 30$ )

$\alpha$	Swap Ops	Total Pages	Final $Th$	Behavior
10	17	68	3	Conservative growth, stable
50	9	147	7	Aggressive growth, bursty I/O
100	5	124	0	Reaches LIMIT immediately

Key observations:

- High  $\alpha$  (50-100): Clears memory quickly but causes I/O spikes
- Moderate  $\alpha$  (20-40): Balances swap size and frequency

#### 0.3.3 Combined Impact Analysis

Table 3: Interaction Effects of  $\alpha$  and  $\beta$

$\alpha$	$\beta$	Swap Ops	Total Pages	Final $Th$	Profile
10	10	31	124	9	Conservative
50	50	6	79	1	Aggressive
30	30	8	74	10	Balanced



### 0.3.4 Optimization Recommendations

- **Stable Workloads:**  $\alpha = 20-40$ ,  $\beta = 30-50$
- **Bursty Workloads:**  $\alpha = 50-70$ ,  $\beta = 10-20$
- **Avoid:**  $\alpha > 80$  (I/O spikes) or  $\beta > 80$  (threshold collapse)

### 0.3.5 Conclusion

The optimal efficiency is achieved with:

- Moderate  $\alpha$  (30-50) to balance swap size and frequency
- Moderate  $\beta$  (20-40) to maintain memory pressure awareness
- Combined tuning using the relationship:

$$\text{Efficiency} \propto \frac{1}{\text{Swap Ops} \times \text{Total Pages}} \quad (3)$$